



Bitmasks

ICS312 Machine-Level and Systems Programming

Henri Casanova (henric@hawaii.edu)

Boolean Bitwise Operations

- There are assembly bitwise instructions for all standard boolean operations: AND, OR, XOR, and NOT
- Bits are computed individually (i.e., “bitwise”)
- Examples:

AND

1	0	1	1	0	0	0	0
1	1	0	1	1	0	1	1
=	1	0	0	1	0	0	0

OR

1	1	0	0	0	1	0	1
0	1	1	0	1	1	1	0
=	1	1	1	0	1	1	1

XOR

1	1	0	0	0	1	0	1
0	1	1	0	1	1	1	0
=	1	0	1	0	1	0	1

NOT

1	1	0	0	0	1	0	1
=	0	0	1	1	1	0	1

Boolean Bitwise Instructions

mov ax, 0C123h

and ax, 082F6h ; ax = C123 **AND** 82F6 = 8022

or ax, 0E34Fh ; ax = 8022 **OR** E34F = E36F

xor ax, 036E9h ; ax = E36F **XOR** 36E9 = D586

not ax ; ax = **NOT** D586 = 2A79

The test Instruction

- The **test** instruction performs an AND, but does not store the result
- It only sets the FLAG bits
 - Just like `cmp` does a subtraction but never stores its result
- Note that all boolean bitwise instructions do set the FLAG bits, **BUT** for the **not** operation, which doesn't
- Example:

```
mov    al, 0FFh
test   al, 000h
jz     foo    ; branches
```

```
mov    al, 0FFh
not    al
jz     foo    ; may not branch
```

Uses of Bitwise operations

- Bitwise operations are used to modify particular bits in data
- This is done via **bit masks**, i.e., constant (or “immediate”) quantities with carefully chosen bits
- Example:
 - Say you want to “turn on” bit 3 of a 2-byte value (counting from the right, with bit 0 being the least significant bit)
 - An easy way to do this is to OR the value with 00000000000001000, which is 8 in decimal
 - Say the value is stored in ax
 - You can simply execute the instruction:
 or ax, 8 ; turns on bit 3 in ax
- Easy to generalize
 - **To turn on bits:** use OR (with appropriate 1’s in the bit mask)
 - **To turn off bits:** use AND (with appropriate 0’s in the bit mask)
 - **To flip bits:** use XOR (with appropriate 1’s in the bit mask)

Bit Mask Operations Examples

mov eax, 04F346BA2h

or ax, 0F000h ; turns on 4 leftmost bits of ax
; eax = 4F34FBA2


xor eax, 000400000h ; flips bit 22 of EAX
; eax = 4F74FBA2

xor ax, 0FFFFh ; 1's complement of ax
; eax = 4F74045D
; (same as doing "not")

Remainder of a Division by 2^i

- To find the remainder of a division of an operand by 2^i , just AND the operand by $2^i - 1$
- Why does this work?

$a = s\text{-bit quantity}$



Therefore, $a = b * 2^i + c$, and c is the remainder!
The remainder is simply the lowest i bits!

Another Way to Look at it

x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

=

x	x	x	x	x	x	x	x	x	x	x	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

+

0	0	0	0	0	0	0	0	0	0	0	x	x	x	x	x
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Example with $i=5$

Another Way to Look at it

a

x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

=

x	x	x	x	x	x	x	x	x	x	x	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

+

0	0	0	0	0	0	0	0	0	0	0	x	x	x	x	x
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Multiple of $2^5 = 32$

Smaller than $2^5=32$

Therefore, $a = \text{<some multiple of 32>} +$
 $\text{<some number smaller than 32>}$

By definition, that smaller than 32 number is the
remainder of the division of a by 32

It is just the low $\log_2(32)=5$ bits of a !!!

Remainder of a Division by 2^i

- Let's compute the remainder of the integer division of 123_{10} by $2^5=32_{10}$ (unsigned) by doing an AND with 2^5-1

mov ax, 123	<table border="1"><tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td><td>1</td><td>0</td><td>1</td><td>1</td></tr></table> ax	0	0	0	0	0	0	0	0	0	0	1	1	1	1	0	1	1
0	0	0	0	0	0	0	0	0	0	1	1	1	1	0	1	1		
mov bx, 0001Fh	<table border="1"><tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr></table> bx	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1
0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1		
and bx, ax	<table border="1"><tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>1</td><td>1</td><td>0</td><td>1</td><td>1</td></tr></table> bx	0	0	0	0	0	0	0	0	0	0	0	0	1	1	0	1	1
0	0	0	0	0	0	0	0	0	0	0	0	1	1	0	1	1		

- The remainder when dividing 123 by 32 is $11011_{10} = 27_{10}$
- Often one says “I **masked out** the 11 high bits of ax” (using bx as the bitmask)

Turning on a specific bit

- Say you want to turn on a specific bit in some data, but that you don't know which one before you run the program
 - the index of the bit to turn on is contained in a register
 - we need to build the bit mask “on the fly”
- Assuming that the index of the bit is initially in `bl`, and that we wish to turn on a bit in `eax`

```
mov  cl, bl          ; must have the bit index in cl
mov  ebx, 1          ; create a number 0...01
shl  ebx, cl         ; shift it left cl times
or   eax, ebx        ; turn on the desired bit using
                      ; ebx as a mask!
```

Turning off a specific bit

- Turning a bit off requires one more instruction, to generate a bit mask that looks like 1...101..1
- Assuming that the index of the bit is initially in bl, and that we wish to turn on a bit in eax

```
mov  cl, bl          ; must have the bit index in cl
mov  ebx, 1          ; create a number 0...01
shl  ebx, cl         ; shift it left cl times
not  ebx;            ; take the complement!
and  eax, ebx        ; turn off the desired bit using
                      ; ebx as a mask!
```

An odd xor

- One often sees the following instruction:
`xor eax, eax ; eax = 0`
- This is a simple way to set `eax` to 0
- It is useful because its machine code is smaller than that of, for instance, “`mov eax, 0`”
- Therefore one saves a few bits in the text segment, meaning that when the program runs it will load fewer bits less from memory, saving perhaps a few cycles
- **Message:** One could do everything with high-level operations, but the good (assembly) programmer (and definitely the good *compiler*) will use bit operations to save memory and/or time
- Let's go through the example in Section 3.3, which is a good example of bit operation “craziness”



Avoiding Conditional Branches

- Section 3.3 is all about a trick to avoid conditional branches
- Conditional branches kill performance
 - Essentially, one key to making processors go fast is to allow them to know what's coming up next
 - With conditional branches, the processor doesn't know in advance whether the branch will be taken or not
- In many cases, one cannot avoid using conditional branches
 - It's just in the nature of the computation
 - For instance, in a for loop
- But in some cases it's possible, with some cleverness
- Let's just look at an example that illustrates some of the coolness/craziness of bitwise operations

SETxx Instructions

- The x86 assembly provides a set of instructions that set the value of a byte register (e.g., al) or of a byte memory location to **0 or 1** based on a flag
- Say you want to set al to 0 if bx > cx or to 1 otherwise (all signed)
- With the **setg** instruction you can save a conditional branch:

; without setg

```
mov  al, 1    ; al = 1
cmp  bx, cx   ; compare
jng  next     ; jump if bx <= cx
mov  al, 0    ; al = 0
next:
```

; with setg

```
cmp  bx, cx
setg al, 0
```

SETxx Instructions

- The x86 assembly provides a set of instructions that set the value of a byte register (e.g., al) or of a byte memory location to **0 or 1** based on a flag
- Set you want to set al to 0 if bx > cx or to 1 otherwise (all signed)
- With the **setg** instruction you can save a conditional branch:

; without setg

```
mov  al, 1    ; al = 1
cmp  bx, cx   ; compare
jng  next     ; jump if bx <= cx
mov  al, 0    ; al = 0
next:
```

; with setg

```
cmp  bx, cx
setg al, 0
```

The second operand is added to al “after” al is set to 0 or 1. In this case we don’t want to add anything, so we use “0”

Example: max(a,b)

- Say we want to store into ecx the maximum of two (signed) numbers, one stored in eax and the other one in [num]
- Here is a simple code to do this

```
cmp    eax, [num]
jge    next          ; conditional branch
mov    ecx, [num]
jmp    end
```

next:

```
mov    ecx, eax
```

end:

- Let's rewrite this without a conditional branch!
- Just follow along... this won't strike you as intuitive but you should be able to understand that it works
 - Even though you wouldn't have been able to do it from scratch on your own, most likely

Example: max(a,b)

- To avoid the conditional branch, one needs a SETxx instruction and clever bit masks

- We use a helper register, ebx, which we set to all zeros

`xor ebx, ebx ; ebx = 0`

- We compare the two numbers

`cmp eax, [num] ; compare eax to [num]`

- We set the value of bl to 0 or 1 depending on the result of the comparison

`setg bl ; set bl to 0 or 1`

□ If `eax > [num]`, `ebx = 1 = 0...01b`

□ If `eax <= [num]`, `ebx = 0 = 0...00b`

- We negate ebx (i.e., take 1's complement and add 1)

`neg ebx`

□ If `eax > [num]`, `ebx = FFFFFFFFh`

□ If `eax <= [num]`, `ebx = 0000000000h`

Example: max(a,b)

- We now have:
 - eax contains one number, [num] contains the other
 - If $\text{eax} > [\text{num}]$, $\text{ebx} = \text{FFFFFFFFh}$ (we want to “return” eax)
 - If $\text{eax} \leq [\text{num}]$, $\text{ebx} = 000000000h$ (we want to “return” [num])
- If eax is the maximum and we AND eax and ebx, we get eax, otherwise we get zero
- If [num] is the maximum and we AND [num] and NOT(ebx), we get [num], otherwise we get zero
- So if we compute $((\text{eax AND ebx}) \text{ OR } ([\text{num}] \text{ AND NOT}(\text{ebx})))$ we get the maximum!
 - If eax is the maximum ($\text{ebx} = \text{FFFFFFFFh}$):
 - $((\text{eax AND ebx}) \text{ OR } ([\text{num}] \text{ AND NOT}(\text{ebx}))) = \text{eax OR } 0\dots 0 = \text{eax}$
 - If [num] is the maximum ($\text{ebx} = 00000000h$):
 - $((\text{eax AND ebx}) \text{ OR } ([\text{num}] \text{ AND NOT}(\text{ebx}))) = 0\dots 0 \text{ OR } [\text{num}] = [\text{num}]$
- Now we just need to compute $((\text{eax AND ebx}) \text{ OR } ([\text{num}] \text{ AND NOT}(\text{ebx})))$

Example: max(a,b)

- Computing ((eax AND ebx) OR ([num] AND NOT(ebx))):

```
mov      ecx, ebx      ;
and      ecx, eax      ; ecx = eax AND ebx
not      ebx           ;
and      ebx, [num]    ; ebx = [num] AND NOT(ebx)
or       ecx, ebx      ; voila!
```


- Whole program:

```
xor      ebx, ebx; ebx = 0
cmp      eax, [num]    ; compare eax and [num]
setg     bl            ; bl = 1 if eax > [num], 0 otherwise
neg      ebx           ; take one's complement + 1
mov      ecx, ebx      ;
and      ecx, eax      ; ecx = eax AND ebx
not      ebx           ;
and      ebx, [num]    ; ebx = [num] AND NOT(ebx)
or       ecx, ebx      ; voila!
```



Predicates

- The technique we have used to remove a conditional branch is often called “predicates”
- Essentially, you replace conditional branches with Boolean computations that “work out”
- Predicates are (or should be) used extensively when programming for the GPU
 - For performance reasons
- Whether the previous example leads to performance gains is unclear:
 - Version #1: we have a “bad” branch instruction
 - Version #2: we have a lot of computation
- But if we have a lot of maximum computations to perform in a row, Version #2 is almost certainly faster



Bitwise Operations in High-Level Languages

- Although in this course we focus on assembly programming, let's discuss high-level languages a little bit
- All high-level languages provide bitwise operations
- Let's look at the typical bitwise operators...

Operators in C/C++/Java/...

- Boolean Operations:

- AND: &&
- OR: ||
- XOR: XXX
- NOT: !

- Bitwise Operations:

- AND: &
- OR: |
- XOR: ^
- NOT: ~

- Shift Operations:

- Left Shift: <<
- Right Shift: >>
- Logical if operand is unsigned (not in Java)
- Arithmetic if operand is signed

Example Operations: C/C++

```
short int          s;    // 2-byte signed
short unsigned int u;    // 2-byte unsigned
s = -1;            // s = 0xFFFF
u = 100;           // u = 0x0064
u = u | 0x0100;    // u = 0x0164
s = s & 0xFFFF0;   // s = 0xFFFF0
s = s ^ u;         // s = 0xFE94
u = u << 3;        // u = 0x0B20
s = s >> 2;        // s = 0xFFA5
```




Common Uses of Bit Operations

- One can use bit operations in high-level languages like in assembly
 - e.g., fast multiplications and divisions
- A very common use is for dealing with file permissions on UNIX systems
- The POSIX API, for dealing with files on all Linux systems, uses bits to encode file access permissions
- If you have to write code that needs to read/modify file permissions, then you need to use bitwise operations

Using chmod from C

- In a POSIX system, there is a C library function called `chmod()` that modifies permissions
- `chmod()` takes two arguments:
 - The file name
 - A 4-byte quantity that is interpreted as a bunch of individual bits, which describe the permission
- To make life easy, the user does not have to construct the bits by hand, but there are macros
- For instance: (p contains the file's permission bits)
`chmod("file", p | S_IRUSR)`
Gives read permission to the owner of the file
S_IRUSR has one of its bits turned on
- This makes it easy to do multiple things at once:
`chmod("file", p | S_IRUSR | S_IWUSR | S_IROTH)`
The user can read and write, all "other" users can read
- Simply use a bitwise OR to apply all permission settings

Modifying Permissions

- Say you want to write a program that, given a file, removes write access to others and adds read access to the owner of the file

- First step: get the 4-byte permission data

```
struct stat  s;      // data structure
stat("file", &s);    // get all file metadata
unsigned int  p;     // 4-byte quantity
p = s.st_mode;       // p = permission bits
```

- Second step: modify, keeping most bits unchanged

```
chmod("file", (p & ~S_IWOTH) | S_IRUSR);
```

Counting Bits

- Section 3.6 of the book shows methods for counting bits
- These methods are shown in C, but of course it's easy (if perhaps cumbersome) to implement them in assembly
- Let's look at Method #1
 - Make sure you look at the others on your own and that you understand them (some are quite creative)

```
unsigned char data; // 1 byte (book uses 4)
char count;        // counter (only 1 byte necessary)
while (data) {
    data = data & (data - 1);
    cnt++;
}
printf("number of 1 bits: %d\n", count);
```

Counting Bits

```
while (data) {  
    data = data & (data - 1);  
    cnt++;  
}
```

■ Example: data = 01011010 (in binary)

- data = data & (data - 1) = 01011010 & 01011001
= 01011000
- data = data & (data - 1) = 01011000 & 01010111
= 01010000
- etc.

- At each step, we set the rightmost 1 bit to 0!
- When we have all zeros we stop
- The number of iterations is the number of 1 bits!!

Detecting duplicates

- Inspired from the first exercise in the “Cracking the Coding Interview” book
- You have write a function that takes an ASCII character as input, and returns true is that function was called previously with that ASCII character
- Typical idea: maintain a set of “already seen” characters
 - e.g., a `<char, boolean>` map, an array of boolean
- But storing a boolean in a byte is super wasteful
 - You’re wasting 7 out of 8 bits
- The book suggests using a “bit field” to save space
- Let’s do it live... (could be bad ;))



Conclusion

- Next week we'll have an in-class practice quiz about this module
- Let's look at Homework assignment #6