

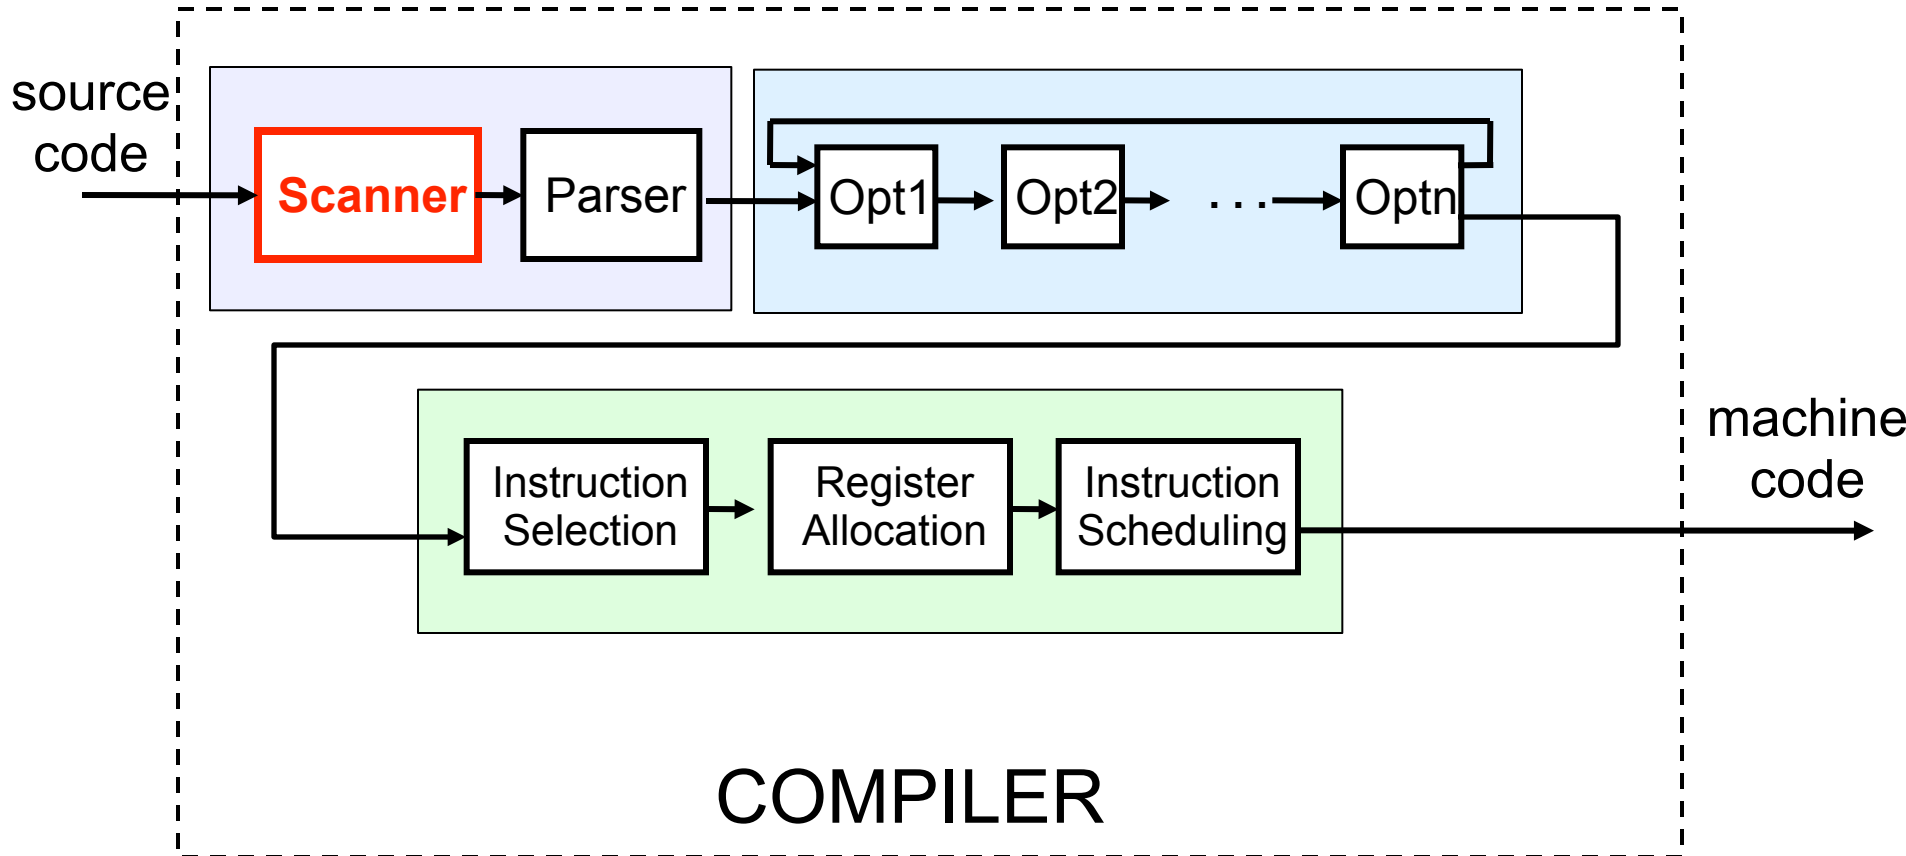


Lexical Analysis (“scanning”, “lexing”)

**ICS312
Machine-Level and
Systems Programming**

Henri Casanova (henric@hawaii.edu)

The Compiler Big Picture



Lexical Analysis

- **Lexical Analysis** is also called ‘scanning’ or ‘lexing’
- It does two things:
 - Transforms the input source string into a sequence of substrings
 - Classifies them according to their “role”
- The input is the source code
- The output is a list of **tokens**
- Example input:

```
if (x == y)
    z = 12;
else
    z = 7;
```

- This is really a single string:

i	f		(x	=	=	y)	\n	\t	z		=		1	2	;	\n	e		s	e	\n	\t	z		=		7	;	\n
---	---	--	---	---	---	---	---	---	----	----	---	--	---	--	---	---	---	----	---	--	---	---	----	----	---	--	---	--	---	---	----

Tokens

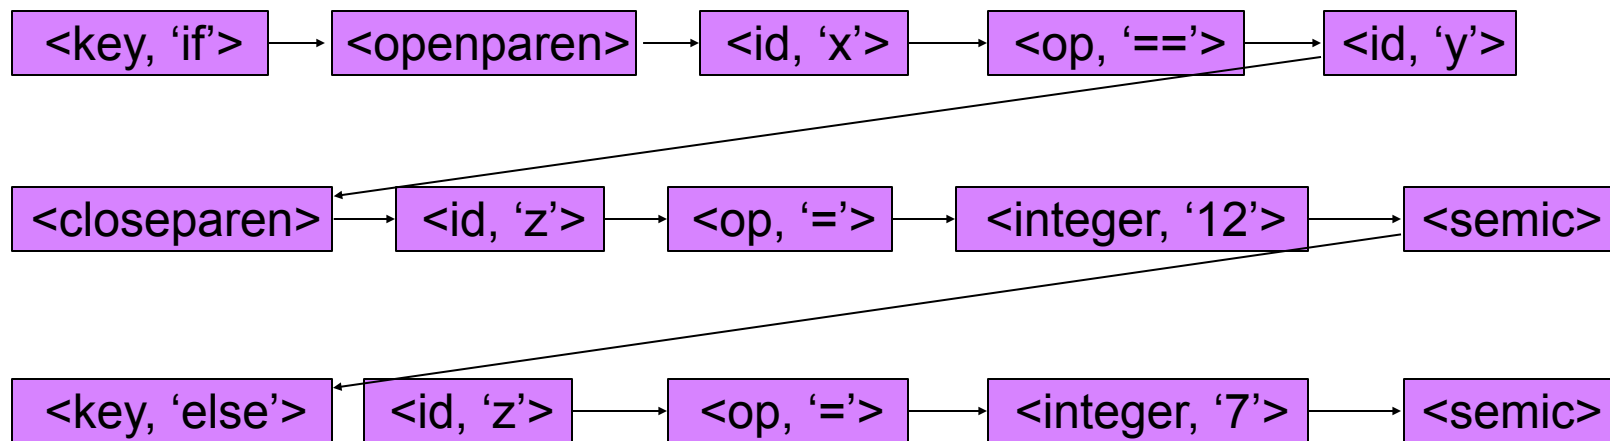
- A token is a **syntactic category**
- Example tokens:
 - Identifier
 - Integer
 - Floating-point number
 - Keyword
 - etc.
- In English we would talk about
 - Noun
 - Verb
 - Adjective
 - etc.

Lexeme

- A **lexeme** is the string that represents an **instance of a token**
 - In English, “big” and “small” would be lexemes off the adjective token
- The set of all possible lexemes that can represent a token instance is described by a **pattern**
 - Good luck with that in English :)
- For instance, we can decide that the pattern for an identifier is
 - A string of letters, numbers, or underscores,

Lexing output

```
i f ( x == y ) \n \t z = 1 2 ; \n e l s e \n \t z = 7 ; \n
```



- Note that the lexer removes non-essential characters
 - Spaces, tabs, linefeeds
 - And comments!
 - Typically a good idea for the lexer to allow arbitrary numbers of white spaces, tabs, and linefeeds

The Lookahead Problem

- Characters are read in from left to right, one at a time, from the input string
- The problem is that it is not always possible to determine whether a token is finished or not without looking at the next character
- Example:
 - Is character 'f' the first letter of a variable name or the first letter of keyword 'for'?
 - Is character '=' an assignment operator or the first character of the '==' operator?
- In some languages, a lot of lookahead is needed
- Example: FORTRAN
 - Fortran removes ALL white spaces before processing the input string (designed before we really understood compilation!)
 - 'DO 5 I = 1.25' is valid code that sets variable DO5I to 1.25
 - But 'DO 5 I = 1.25' could also be the beginning of a for loop!

The Lookahead Problem (2)

- It is typically a good idea to design languages that require ‘little’ lookahead
 - For each language, it should be possible to determine how many lookahead characters are needed
- Example with 1-character lookahead:
 - Say that I get an ‘if’ so far
 - I can *look* at the next character
 - If it’s a ‘ ‘, ‘(’, ‘\t’, then I don’t read it; I stop here and emit a TOKEN_IF
 - Otherwise I read the next character and will most likely emit a TOKEN_ID
- In practice one implements lookahead/pushback
 - When in need to look at next characters, read them in and push them onto a data structure (stack/fifo)
 - When in need of a character get it from the data structure, and if empty from the file

A Lexer by Hand?

- Example: Say we want to write the code to recognizes the keyword 'if'

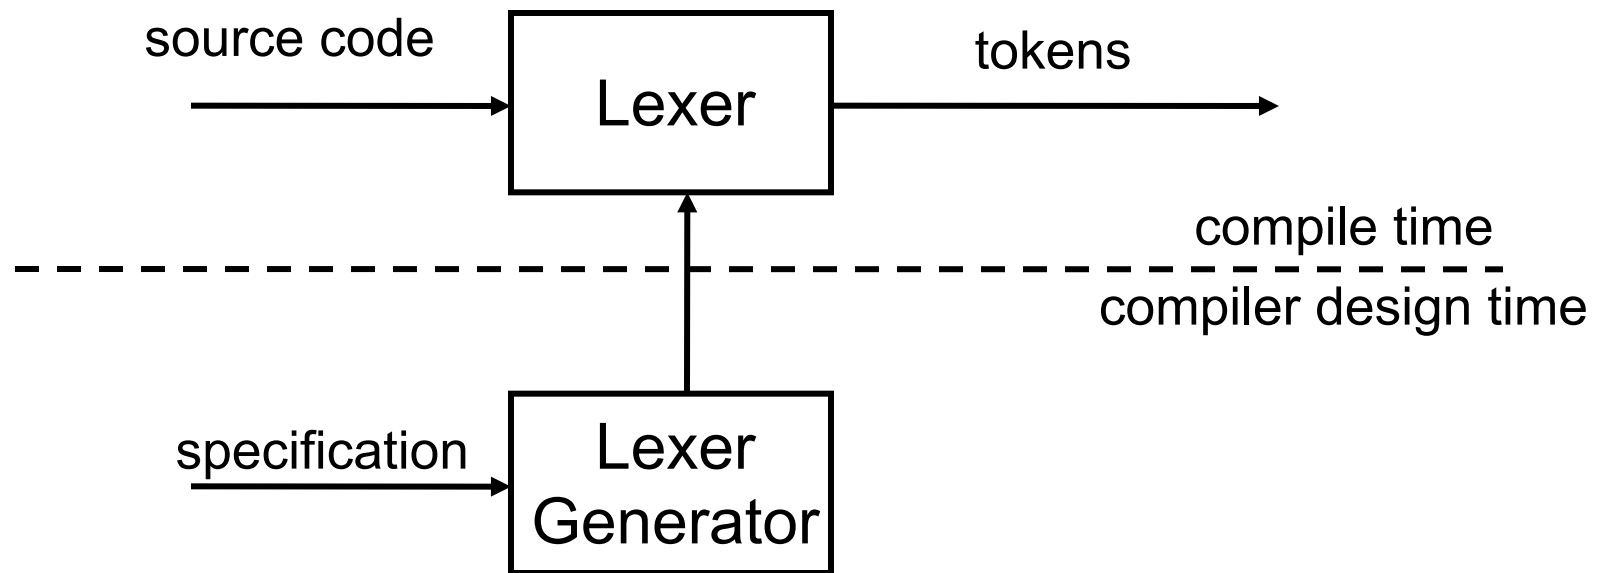
```
c = readchar();
if (c == 'i') {
    c = readchar();
    if (c == 'f') {
        c = readchar();
        if (c not alphanumeric) {
            pushback(c);
            emit(TOKEN_IF)
        } else {
            // build a TOKEN_ID
        }
    } else {
        // something else
    }
} else {
    // something else
}
```

A Lexer by Hand? You're joking!

- There are many difficulties when writing a lexer by hand as in the previous slide
 - Many types of tokens
 - fixed string
 - special character sequences (operators)
 - numbers defined by specific/complex rules
 - Many possibilities of token overlap
 - Hence, many nested if-then-else in the lexer's code
- Coding all this by hand is very, very painful
 - And it's almost impossible to get it right
 - Nevertheless, some compilers have an implemented-by-hand lexer for higher speed!

Automatic Lexer Generation?

- To avoid the endless nesting of if-then-else one needs a formalization of the lexing process
- If we have a good formalization, we could even **generate the lexer's code automatically!**



Lexer Specification

- Question: How do we tell the lexer how to recognize the tokens of a specific language?
- Essentially we need a way to describe each token accurately and completely
- Example: “an integer is a sequence of 1 or more digits, where a digit is 0, 1, 2, 3, 4, 5, 6, 7, 8, or 9”
- Example: “an assignment operator is an equal sign”
- Example: “a equality operator is a sequence of two equal signs”
- Example: “a variable name is a sequence of any alphanumeric character, but it can’t start with a digit, and it can’t have = or ; in it, but it’s ok if it has _ or,....”
- Example: “a floating point number is.... well, either a + or - or nothing, and then an integral part or not, and a decimal point, or not... wait...also there can be an exponent, with is denoted by ‘e’ or ‘E”, followed by an integer unless... okokok, let me start again....”
- Doing this in plain English really gets difficult!

Lexer Specification

- We need to formalize the process in the previous slide
- How do we formalize things? We define a language!
- What's a language?
 - An alphabet (typically denoted by Σ)
 - in our case: the ASCII characters allowed in source code
 - A subset of all the possible strings over Σ
- We need to provide a formal definition of the language of the tokens over Σ
 - Define which strings are valid tokens
 - All other strings are invalid
- It turns out that for all (reasonable) programming languages, the tokens can be described by **regular expressions**
 - See ICS 222, ICS 241, ICS 313

Regular Expressions

- Regular expressions are notations

- A regular expression is a string (in a meta-language) that describes a pattern (in the token language)
- You have seen them in ICS 241!

Expression	Meaning
ϵ	empty pattern
a	One occurrence of pattern 'a'
ab	Strings with pattern 'a' followed by pattern 'b'
$a b$	Strings with pattern 'a' or pattern 'b'
a^*	Zero or more occurrences of pattern 'a'
a^+	One or more occurrences of pattern 'a'
$a?$	$(a \epsilon)$
$.$	Any single character (not very standard)

REs for Keywords

- It is easy to define a RE that describes all keywords

Keyword = 'if' | 'else' | 'for' | 'while' | 'int' | ..

- These can be split in groups if needed

Keyword = 'if' | 'else' | 'for' | ...

Type = 'int' | 'double' | 'long' | ...

- The choice depends on what the next component (i.e., the parser) would like to see

RE for Numbers

- Straightforward representation for integers
 - `digit = '0' | '1' | '2' | '3' | '4' | '5' | '6' | '7' | '8' | '9'`
 - `integer = digit+`
- RE systems allow ranges, sometimes with '[' and ']'
 - `integer = [0-9]+`
- Floating point numbers are much more complicated
 - `2.00, .12e-12, 312.00001E+12, 4, 3.141e-12`
- Here is one attempt
 - `('+'|'| '-'|ε)(digit+ '.'? | digits* ('.' digit+)) (('E'|'e')('+'|'| '-'|ε) digit+))?`
- Note the difference between meta-character and language-characters
 - '+' versus +, '-' versus -, '(' versus (, etc.
- Often books/documentations use different fonts for each level of language
 - In these slides I use single quotes for the “programming” language

RE for Identifiers

- Here is a typical description
 - letter = 'a'-'z' | 'A'-'Z'
 - ident = letter (letter | digit | '_')^{*}
 - Starts with a letter
 - Has any number of letter or digit or '_' afterwards
- In C: ident = (letter | '_') (letter | digit | '_')^{*}
- Let's try to compile "int 3x;" with gcc..

RE for Phone Numbers

- Simple RE

- `digit = ['0'-'9']`
- `area = digit digit digit`
- `exchange = digit digit digit`
- `local = digit digit digit digit`
- `phonenumber = '(' area ')' ' '? exchange ('-'|') local`

- The above describes the 10^{3+3+4} strings of the $L(\text{phonenumber})$ language

REs in Practice

- The Linux grep utility allows the use of REs
 - Example with phone numbers
 - `grep '([0-9]\{3\}) \{0,1\}[0-9]\{3\}[-|][0-9]\{4\}' file`
 - The syntax is different from that we've seen, but equivalent
 - Sadly, there is no single standard for RE syntax
- Perl/Python/Java/... all implement regular expressions, all with similar-ish syntaxes
- (Good) text editors implement regular expressions
 - .e.g., for string replacements
 - Let's see it for vim and sed
- At the end of the day, we often use tons of regular expressions
- And most programs you use everyday use REs internally
- And of course, so do compilers

In-class Exercise #1

- Which of the regular expressions below corresponds to: “All non-empty strings over alphabet {a,b,c}”? (poll)
 - ☐ **#1:** $a^* b^* c^*$
 - ☐ **#2:** $a^+ b^+ c^+$
 - ☐ **#3:** $(a \mid b \mid c)^+$
 - ☐ **#4:** $a^+ \mid b^+ \mid c^+$
 - ☐ **#5:** $a \mid b \mid c$

In-class Exercise #1 (solution)

- Which of the regular expressions below corresponds to: “All non-empty strings over alphabet {a,b,c}”?
 - ☐ #1: $a^* b^* c^*$
 - ☐ #2: $a^+ b^+ c^+$
 - ☐ #3: $(a \mid b \mid c)^+$
 - ☐ #4: $a^+ \mid b^+ \mid c^+$
 - ☐ #5: $a \mid b \mid c$

In-class Exercise #2

- Write a regular expression for “All strings over alphabet {a,b,c} that contain substring ‘abc’”?

In-class Exercise #2 (solution)

- Write a regular expression for “All strings over alphabet {a,b,c} that contain substring ‘abc’”?

$(a \mid b \mid c)^* a b c (a \mid b \mid c)^*$

In-class Exercise #3

- Write a regular expression for “All strings over alphabet $\{a,b,c\}$ that consist of one or more a’s, followed by two b’s, followed by whatever string of a’s and c’s”?

In-class Exercise #3 (solution)

- Write a regular expression for “All strings over alphabet {a,b,c} that consist of one or more a’s, followed by two b’s, followed by whatever string of a’s and c’s”?

$a^+ b b (a | c)^*$

In-class Exercise #4

- Which of the REs below correspond to “All strings over alphabet $\{a,b,c\}$ such that they contain at least one of substrings ‘abc’ or ‘cba’”?
- **#1:** $(abc)^+ \mid (cba)^+$
- **#2:** $(a \mid b \mid c)^* (a \mid c) b (a \mid c) (a \mid b \mid c)^*$
- **#3:** $(a \mid b \mid c)^* (a b c) (c b a) (a \mid b \mid c)^*$
- **#4:** $(a \mid b \mid c)^* ((a b c) \mid (c b a)) (a \mid b \mid c)^*$
- **#5:** $(a^* b^* c^* a b c a^* b^* c^*) \mid (a^* b^* c^* c b a^* b^* c^*)$
- **#6:** $((a \mid b \mid c)^* a b c (a \mid b \mid c)^*) \mid ((a \mid b \mid c)^* c b a (a \mid b \mid c)^*)$

In-class Exercise #4 (solution)

- Which of the REs below correspond to “All strings over alphabet $\{a,b,c\}$ such that they contain at least one of substrings ‘abc’ or ‘cba’”?
- #1: $(abc)^+ \mid (cba)^+$
- #2: $(a \mid b \mid c)^* (a \mid c) b (a \mid c) (a \mid b \mid c)^*$
- #3: $(a \mid b \mid c)^* (a b c) (c b a) (a \mid b \mid c)^*$
- #4: $(a \mid b \mid c)^* ((a b c) \mid (c b a)) (a \mid b \mid c)^*$
- #5: $(a^* b^* c^* a b c a^* b^* c^*) \mid (a^* b^* c^* c b a^* b^* c^*)$
- #6: $((a \mid b \mid c)^* a b c (a \mid b \mid c)^*) \mid ((a \mid b \mid c)^* c b a (a \mid b \mid c)^*)$


Now What?

- Now we have a nice way to formalize each token (which is a set of possible strings)
- Each token is described by a RE
 - And hopefully we have made sure that our REs are correct
 - Easier than writing the lexer from scratch!
 - But still requires that one be careful
- **Question:** How do we use these REs to parse the input source code and generate the token stream?
- A little bit of theory:
 - REs characterize Regular Languages
 - Regular Languages are recognized by Finite Automata (which you have also seen in ICS241)


Finite Automata

- A finite automaton is defined by
 - An input alphabet: Σ
 - A set of states: S
 - A start state: n
 - A set of accepting states: F (a subset of S)
 - A set of transitions between states: subset of $S \times S$
- Transition Example
 - $s_1: a \rightarrow s_2$
 - If the automaton is in state s_1 , reading a character 'a' in the input takes the automaton in state s_2
 - Whenever reaching the 'end of the input,' if the state the automaton is in is an accept state, then we accept the input
 - Otherwise we reject the input


Finite Automata as Graphs




■ A state



■ A start state




■ An accepting state




■ A transition

Automaton Examples



- This automaton accepts input 'if'

Automaton Examples




- This automaton accepts strings that start with a 0, then have any number of 1's, and end with a 0
- Note the natural correspondence between automata and REs:
 01^*0
- **Question:** can we represent all REs with finite automata?
- **Answer:** yes
- Therefore, if we write a piece of code that implements arbitrary automata, we have a piece of code that implements arbitrary REs, and we have a lexer!
 - Not this simple, but really close

Non-deterministic Automata

- The automata in the previous slide are called **Deterministic Finite Automata (DFA)**
 - At each state, there is at most one edge for a given symbol
 - At each state, transition can happen only if an input symbol is read
 - Or the string is rejected
- It turns out that it's easier to translate REs to Non-deterministic Finite Automata (NFA)
 - There can be ' ϵ -transitions'
 - Taken arbitrarily without consuming an input character
 - There can be multiple possible transitions for a given input symbol at a state
 - The automaton can take them all simultaneously (see later)


Example REs and DFA

- Say we want to represent RE 'a*b*c*d*e' with a DFA



Example REs and NFA


- 'a*b*c*d*e': much simpler with a NFA



- With ϵ -transitions, the automaton can 'choose' to skip ahead, non-deterministically
- This means that the automaton keeps track of all possible paths the automaton could have taken based on its input

Example REs and NFA


- 'a+b+c+d+e': easy modification



- But now we have multiple choices for a given character at each state!
 - e.g., two 'a' arrows leaving n

NFA Acceptance


- When using an NFA, one must constantly keep track of all possible states
- If at the end of the input (at least) one of these states is an accepting state, then accept, otherwise reject



input string: 010

NFA Acceptance


- When using an NFA, one must constantly keep track of all possible states
- If at the end of the input (at least) one of these states is an accepting state, then accept, otherwise reject



input string: 010 (after reading a 0 I could be in any state!)

NFA Acceptance


- When using an NFA, one must constantly keep track of all possible states
- If at the end of the input (at least) one of these states is an accepting state, then accept, otherwise reject



input string: 0**1**0 (after reading a 1, I can only be in state n)

NFA Acceptance


- When using an NFA, one must constantly keep track of all possible states
- If at the end of the input (at least) one of these states is an accepting state, then accept, otherwise reject



input string: 010 (after reading a 0, I can be in any state)

NFA Acceptance

- When using an NFA, one must constantly keep track of all possible states
- If at the end of the input (at least) one of these states is an accepting state, then accept, otherwise reject




input string: 010

ACCEPT because we **could** be in s2 after the input has been processed entirely

In-class exercise#5


- Which RE corresponds to this automaton



- #1: $a a^* b a$
- #2: $a^+ b^+ a$
- #3: $a a^* b^* a$
- #4: $a^* b^* a$
- #5: $a^+ b^*$

In-class exercise#5 (solution)


- Write a RE for this automaton



- #1: $a a^* b a$
- #2: $a^+ b^+ a$
- #3: $a a^* b^* a$ (equivalent to $a^+ b^* a$)
- #4: $a^* b^* a$
- #5: $a^+ b^*$


In-class exercise #6

- Write a RE for this NFA



In-class exercise #6 (Solution)


- Write a RE for this NFA



$a^* b^* \mid a^* b^* a b^*$

In-class exercise #7


- Which of the REs below are for this automaton?
(poll)



- #1: $a^* b^* a b^* a$ | $a^* b^+ a b^+$
- #2: $a^* b^* a b^* a$ | $a^* b^+ a b^*$
- #3: $a^* b^+ a b^* a$ | $a^* b^+ a b^+$
- #4: $a^+ b^* a$ | $a^* b^+ a b^+$

In-class exercise #7 (solution)


- Which of the REs below are for this automaton?



- #1: $a^* b^* a b^* a$ | $a^* b^+ a b^+$
- #2: $a^* b^* a b^* a$ | $a^* b^+ a b^*$
- #3: $a^* b^+ a b^* a$ | $a^* b^+ a b^+$
- #4: $a^+ b^* a$ | $a^* b^+ a b^+$

In-class exercise #8


- Which of the REs below are for this automaton?



- #1: $a^* b c^+ (b^* \mid b^* a \mid c b^* a^*)$
- #2: $a^* b c^* (b^* a \mid c b^* a^*)$
- #3: $a^* b c^* (b^* a? \mid c (b a)^*)$
- #4: $a^* b c^* (b^* a \mid c (b a)^*)$

In-class exercise #8 (solution)


- Which of the REs below are for this automaton?



- #1: $a^* b c^+ (b^* \mid b^* a \mid c b^* a^*)$
- #2: $a^* b c^* (b^* a \mid c b^* a^*)$
- #3: $a^* b c^* (b^* a? \mid c (b a)^*)$
- #4: $a^* b c^* (b^* a \mid c (b a)^*)$


In-class exercise #9

- Write an RE for this automaton



In-class exercise #9 (solution)

- Write an RE for this automaton



$(a^* b c)^* b c (ba)^* b?$


REs and NFA

- So now we're left with two possibilities
- Possibility #1: design DFAs
 - Easy to follow transitions once implemented
 - But really cumbersome
- Possibility #2: design NFAs
 - Really trivial to implement REs as NFAs
 - But what happens on input characters?
 - Non-deterministic transitions
 - Should keep track of all possible states at a given point in the input!
- It turns out that:
 - NFAs are not more powerful than DFAs
 - There are systematic algorithms to convert NFAs into DFAs and to limit their sizes
 - There are simple techniques to implement DFAs in software quickly

Putting it All Together

- Steps to designing/building a lexer
 - Come up with a RE for each token category
 - Come up with an NFA for each RE (easy for a human to do)
 - Convert the NFA (automatically) to a DFA (we have great algorithms to do it)
 - Write a piece of code that implements a DFA (easy for a human to do with a simple transition table)
 - Implement your lexer as a ‘bunch of DFAs’
- Let’s see an example of DFA

Example DFA Implementation



state	char	next state	decision / continue
n	0	s1	REJECT / YES
n	1	-	REJECT / NO
s1	0	s2	ACCEPT / YES
s1	1	s1	REJECT / YES
s2	0	-	REJECT / NO
s2	1	-	REJECT / NO

```
state = STATE_N
while (c == get_next_character()) {
    transition(state,c,&next_state,
    &decision, &continue);
    if (!continue)
        break;
    state = next_state
}
return decision;
```

The ‘bunch of automata’

- How the lexer works
 - The lexer is simply a ‘bunch of automata’
 - It runs them all at the same time until they have all rejected the input
 - It then rewinds to the one that accepted last
 - this is the one that accepted the longest string
 - ‘rewinding’ uses lookahead/pushback
 - This one corresponds to the right token
- Let’s look at this on an example


Example

- Say we have the following tokens (described by a RE, and thus an NFA, and thus a DFA):
 - TOKEN_IF: 'if'
 - TOKEN_IDENT: letter (letter | '_')+
 - TOKEN_NUMBER: (digit)+
 - TOKEN_COMPARE: '=='
 - TOKEN_ASSIGN: '='
- This is a very small set of tokens for a tiny language
- The language assumes that tokens are all separated by spaces
- Let's see what happens on the following input:

i	f		i	f	0		=	=		c		x		=		2	x	3
---	---	--	---	---	---	--	---	---	--	---	--	---	--	---	--	---	---	---

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	ok so far
TOKEN_IDENT	ok so far
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	ok so far
TOKEN_IDENT	ok so far
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Example

i f i f 0 = = c x = 2 x 3




DFA	DECISION
TOKEN_IF	accept
TOKEN_IDENT	accept
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Both TOKEN_IF and
TOKEN_IDENT were the
last ones to accept

Emit **TOKEN_IF** because
we build our lexer with the
notion of **reserved**
keywords

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	ok so far
TOKEN_IDENT	ok so far
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	ok so far
TOKEN_IDENT	ok so far
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Example


i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	ok so far
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Example

i f 0 = = c x = 2 x 3




DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	accept
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Emit **TOKEN_IDENT**
(with string 'if0') because
it accepted the latest

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	reject
TOKEN_NUMBER	reject
TOKEN_COMPARE	ok so far
TOKEN_ASSIGN	ok so far

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	reject
TOKEN_NUMBER	reject
TOKEN_COMPARE	ok so far
TOKEN_ASSIGN	reject

Example

i f i f 0 = = c x = 2 x 3




DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	reject
TOKEN_NUMBER	reject
TOKEN_COMPARE	accept
TOKEN_ASSIGN	reject

Emit **TOKEN_COMPARE**
because it accepted the
latest

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	ok so far
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Example

i f i f 0 = = c x = 2 x 3




DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	accept
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Emit **TOKEN_IDENT**
(with string 'c') because it
accepted the latest

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	ok so far
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Example

i f i f 0 = = c x = 2 x 3




DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	accept
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Emit **TOKEN_IDENT**
(with string 'x') because it
accepted the latest

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	reject
TOKEN_NUMBER	reject
TOKEN_COMPARE	ok so far
TOKEN_ASSIGN	ok so far

Example

i f i f 0 = = c x = 2 x 3




DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	reject
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	accept

Emit **TOKEN_ASSIGN**
because it was the only
one accepted

Example


i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	reject
TOKEN_NUMBER	ok so far
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject

Example

i f i f 0 = = c x = 2 x 3



DFA	DECISION
TOKEN_IF	reject
TOKEN_IDENT	reject
TOKEN_NUMBER	reject
TOKEN_COMPARE	reject
TOKEN_ASSIGN	reject


Abort and print a
Syntax Error
Message!!

Example

- If there had been no syntax error, the lexer would have emitted:
 - <TOKEN_IF>
 - <TOKEN_ID, 'if0'>
 - <TOKEN_COMPARE>
 - <TOKEN_ID, 'c'>
 - <TOKEN_ID, 'x'>
 - <TOKEN_ASSIGN>
 - <TOKEN_NUMBER, '23'>


Implementing 'bunch of DFAs'

- We have one NFA per token
- We can easily combine them in one single NFA




Implementing 'bunch of DFAs'

- We have one NFA per token
- We can easily combine them in one single NFA



Implementing 'bunch of DFAs'

- We have one NFA per token
- We can easily combine them in one single NFA



- We can then convert it to a (massive!) DFA

Implementing a Lexer

- Implementing a Lexer is now straightforward
 - Come up with a RE for each token category
 - Come up with an NFA for each RE
 - Merge them into a bunch of NFAs
 - Convert this (huge) NFA (automatically) to a (enormous) DFA
 - Write a piece of code that implements a DFA
 - And voila!
- The above has been understood for decades
- And it's so “easy” that we now have automatic lexer generators!
- Well-known examples are **lex** and **flex**
- Let's look at another one: **ANTLR**



ANTLR

- ANTLR: A tool to generate lexer/parsers
- It is extremely popular because very simple and yet very powerful
- Let's look on the course Web site at the “Installing ANTLR” page

ANTLR: let's build a lexer

- Say we want to define a language with the following:
 - Reserved keywords: int, if, endif, print
 - An addition operator: '+'
 - An assignment operator: '='
 - An equal operator: '=='
 - A not-equal operator: '!='
 - Integers
 - Variable names as strings of lower-case letters
 - Semicolons for terminating statements
 - Left and right parentheses
 - The ability to ignore white spaces, tabs, carriage returns, etc.

ANTLR: let's build a lexer

- Basics of Regular Expressions in ANTLR:
 - Regular expression name (chosen by you)
 - Colon
 - Regular expression
 - Semicolon
- Example:
 - DIGIT : [0-9] ;
 - VARIABLE: [a-z]+ ;
 - EQUAL: '==' ;
- Let's look at the “A Simple ANTLR Lexer” page on the Web site, and run it...
 - Not that this example has some “parser stuff” at the beginning, but we're ignoring that for now

Conclusion

- 20,000 ft view
 - Lexing relies on Regular Expressions, which rely on NFAs, which are easy to design, which are equivalent to DFAs, which are easy to implement
 - Therefore lexing is “easy” :)
- Lexing has been well-understood for decades and lexer generators are known
 - We’ve seen and will use ANTLR
- The only motivation to write a lexer by hand today: speed