



Arithmetic

ICS312 Machine-Level and Systems Programming

Henri Casanova (henric@hawaii.edu)

Addition and Subtraction

- We've seen two instructions for additions and subtractions: **add** and **sub**
- Both instructions can be used on a pair of signed numbers or on a pair of unsigned numbers
 - One of the big advantages of 2's complement storage
 - **No mixing of signed and unsigned numbers**
- **IMPORTANT (AGAIN):** The CPU does not know whether numbers stored in registers are signed or unsigned!
 - You, the programmer, must keep your own interpretation of the number consistent throughout your program
 - The CPU will happily add whatever registers together using binary addition
- Let's see an example that shows that addition works on either signed or unsigned numbers

The Magic of 2's Complement

- I have two 1-byte values, A3 and 17, and I add them together:

$$A3 + 17 = BA$$

- If my interpretation of the numbers is **unsigned**:

- A3h = 163d
- 17h = 23d
- BAh = 186d
- and indeed, $163d + 23d = 186d$

- If my interpretation of the numbers is **signed**:

- A3h = -93d
- 17h = 23d
- BAh = -70d
- and indeed, $-93d + 23d = -70d$

- So, as long as I stick to my interpretation, the binary addition will do the right thing.... amazing!

- Same thing for the subtraction

Overflow :(

- The magic works only with numbers of within acceptable ranges
- Because we encode numbers with finite numbers of bits, sometimes we want to store a number that would require more bits
 - The results is out of range
- In this case we have “overflow”
- The CPU proceeds with the computations, but “drops” bits that can’t fit
 - Again, it has no idea what the numbers mean
- The numerical result of the operations is then wrong and your program buggy

Overflow and Range (1-byte)

- 1-byte unsigned numbers have range 0, 255
- 1-byte signed numbers have range -128, + 127
- Example additions
 - adding 1-byte unsigned quantity 240d to 1-byte unsigned quantity 100d will lead to an overflow because $340d > 255d$
 - subtracting 1-byte unsigned quantity 240d from 1-byte unsigned quantity 100d will lead to an overflow because $-140d < 0d$
 - adding 1-byte signed quantity 100d to 1-byte signed quantity 120d will lead to an overflow because $220d > 127d$
 - etc.
- Let's see how, as humans, we can detect/understand overflow...
 - Of course one full-proof way is to convert everything to decimal and check whether the result is in range
 - But often we can simply reason about the numbers..



Unsigned Overflow

- Say all our numbers are meant to be unsigned
 - Again, the CPU has no idea about signed/unsigned
- We have overflow when:
 - An addition would lead to a result that can't be encoded in the required number of bits
 - i.e., add something big to something big
 - i.e., we have a left-over carry that can't fit in the result
 - A subtraction would lead to a negative result
 - i.e., subtract something big from something small
- Let's see 1-byte examples...

Unsigned Overflow Examples

■ 1-byte Example (all in hex):

- FF + 02 OVERFLOW (result would be 101h)
 - $255 + 2 > 255$
- 01 - 05 OVERFLOW (result cannot be negative)
 - $1 - 5 < 0$
- 8A - 0F NO OVERFLOW (result is 7Bh)
 - $138 - 15 = 123$
 - We're subtracting something small (0F) from something big (8A), so we can't be negative

■ In a nutshell

- An addition/sub overflows if there is a leftover carry
- BIGGER - SMALLER never overflows
- SMALLER - BIGGER always overflows

In-Class Exercise: Unsigned

■ Which of these **unsigned** operations cause overflow?

- ☐ $0F12 + F212$ (2-byte quantities)
- ☐ $00E3 + F74F$ (2-byte quantities)
- ☐ $F1 - FA$ (1-byte quantities)
- ☐ $FB12 - A3AA$ (2-byte quantities)
- ☐ $A314 - B010$ (2-byte quantities)

In-Class Exercise: Solutions

- Which of these unsigned operations cause overflow?

0F12
+ F212
= 10124 OVERFLOW

00E3
+ F74F
= F832 NO OVERFLOW

- F1 - FA: smaller - bigger OVERFLOW
- FB12 - A3AA: bigger - smaller NO
- OVERFLOW
- A314 - B010: smaller - bigger OVERFLOW

Signed Overflow

- It's a bit more difficult to think about the range for signed numbers because both positive and negative values are possible
- 1-byte Example (all in hex, same as before):
 - $FF + 02$ NO OVERFLOW (result is 01h)
 - $-1 + 2 = +1$
 - $01 - 02$ NO OVERFLOW (result is FFh)
 - $1 - 2 = -1$
 - $8A - 0F$ OVERFLOW (result would be $< 80h$)
 - 8A is negative, and is equal to $-76h = -118d$
 - $-118 - 15 < -128$, and thus cannot be represented as a 1-byte signed quantity
 - We subtracted from a number that's very close to the edge of the valid range
- **Useful trick:** look whether the sign of the result is in agreement with the sign of the operands (next slide)

Signed Overflow

- A possible way to determine whether a particular signed operation would overflow is to look at the sign of the result and see if it makes sense
- Same example as before: $8A - 0F$
 - $8A < 0$ and $0F > 0$, so the result should be negative
 - Lets compute the result
 - I don't like to hex subtractions, so I instead compute $-0F = +F1$
 - “flip and add one” to get the opposite number
 - In hex $8A + F1 = 7B$ (we dropped the carry to fit in 8 bits)
 - $7B$ is positive! OVERFLOW
- In a nutshell:
 - POSITIVE + POSITIVE should be POSITIVE
 - NEGATIVE + NEGATIVE should be NEGATIVE
 - POSITIVE + NEGATIVE never causes overflow!

In-Class Exercise: Signed

- Which of these **signed** operations cause overflow?
 - ☐ $00E3 + FF4F$ (2-byte quantities)
 - ☐ $F1 - 7A$ (1-byte quantities)
 - ☐ $FF847CAA + 78AA0401$ (4-byte quantities)
 - ☐ $DF + EF$ (1-byte quantities)

In-Class Exercise: Solutions

- Which of these signed operations cause overflow?
 - $00E3 + FF4F$
 - POSITIVE + NEGATIVE: NO OVERFLOW
 - $F1 - 7A$
 - I do the hex addition: $F1 - 7A = F1 + 86 = 77$
 - Should be negative: OVERFLOW
 - $FF847CAA + 78AA0401$
 - NEGATIVE + POSITIVE: NO OVERFLOW
 - $DF + EF$
 - SMALL NEGATIVE + SMALL NEGATIVE: NO OVERFLOW
 - $DF + EF = CD$ (dropped a carry), which is negative

Unsigned Overflow

On web site as
ics312_overflow_unsigned.asm

mov	al, 0F0h	; al = F0h
mov	bl, 0A3h	; bl = A3h
add	al, bl	; al = al + bl
movzx	eax, al	; increase size for printing
call	print_int	; print al as an integer

- As a programmer we decided to do some computation with **unsigned values**
- We put value F0h in al (unsigned F0h is decimal 240)
- We put value A3h in bl (unsigned A3h is decimal 163)
- We add them together
- The “true” result should be decimal $240 + 163 = 403$, which cannot be encoded on 8 bits (should be < 255)
- But the processor just goes ahead: $F0 + A3 = 193h$, and then drops the leftmost bits to truncate to a 1-byte value to get 93h!
- To call print_int, we need the integer in eax, so we movzx al into eax
- print_int prints the decimal value corresponding to 00000093h, that is: 147!
- This is obviously wrong, and we can tell (or will be able to shortly) because the carry bit is in fact set to 1
- **Note that this is all correct if we assume signed values and replace movzx by movsx, but then our initial interpretation of the two values is different**

Signed Overflow

On web site as
ics312_overflow_signed.asm

mov	al, 09Ah	; al = 9Ah
mov	bl, 073h	; bl = 73h
sub	al, bl	; al = al - bl
movsx	eax, al	; increase size for printing
call	print_int	; print al as an integer

- As a programmer we decided to do some computation with **signed values**
- We put value 9Ah in al (signed 9Ah is decimal -102)
- We put value 73h in bl (signed 73h is decimal +115)
- We subtract bl from al
- The “true” result should be decimal $-102 - 115 = -217$, which cannot be encoded on 8 bits (should be ≥ -128)
- But the processor just goes ahead: $9Ah - 73h = 9Ah + 8Dh = 27h$
- To call print_int, we need the integer in eax, so we movsx al into eax
- print_int prints the decimal value corresponding to 00000027h, that is: 39!
- This is obviously wrong, and we can tell (or will be able to shortly) because the overflow bit is in fact set to 1
- **Note that this is all correct if we assume unsigned values and replace movsx by movzx, but then our initial interpretation of the two values is different**



Overflow is your Responsibility

- The processor merely computes bits and puts them into the destination location, possibly dropping bits, as if everything were fine, and it's your responsibility to check the overflow!
- In your program you should have checks for overflow, which is annoying
 - That's true in high-level languages as well!
 - Which is why we often use too many bits (e.g., 4-byte values for numbers we know to be small)
 - This way we waste memory, but we're pretty sure to avoid overflow
 - Until we don't all everything falls apart

The FLAG register

- You probably have forgotten by now, but at the beginning of the semester I mentioned the FLAG register
- It's basically a bunch of bits that are set/unset when instructions are executed
- They have a range of uses
- Two of them have to do with overflow:
 - The carry bit
 - The overflow bit

Detecting Overflow in Code

- There is an overflow with an unsigned operation (i.e., on unsigned quantities) if the **carry bit in the FLAG register is set**
 - If the carry bit is set, that means we'd need a larger quantity to hold the result
 - This also works for subtractions (instead of a carry, we have a "borrow", but it's still set in the carry bit)
- There is an overflow with a signed operation (i.e., on signed quantities) if the **overflow bit in the FLAG register is set**
 - This bit is set when the sign of the result does not agree with the signs of the operands
- We'll see later how to check those bits

To remember

domain	overflow
unsigned	carry bit
signed	overflow bit

- After a valid unsigned operation, the overflow bit could be set
- After a valid signed operation, the carry bit could be set

Multiplication

- There are two instructions to perform multiplications
- Multiplying **unsigned** numbers: **mul**
- Multiplying **signed** numbers: **imul**
- Why do we need two different instructions?
- Consider the multiplication of FF by FF
 - If we assume unsigned quantities, this is $255 * 255 = 65035 = \text{FE0Bh}$
 - If we assume signed quantities, this is $-1 * -1 = 1 = 0001\text{h}$
- So clearly we need two different instructions because we need to get two different results

The mul Instruction

- The size of the result of the multiplication is sometimes twice larger than the size of the operands
 - Multiplications just leads to much bigger numbers than additions
 - At most the result will be twice the size of the operands ($255 * 255 = 65,025$, which is encodable on 2 bytes)
- The oldest form of multiplication is the “mul” instruction, which produces a result twice the size of its **unsigned** operand

mul <register or memory reference>

 - If the operand is a byte, then it is multiplied by AL and the result is stored in (16-bit) AX
 - If the operand is 16-bit, it is multiplied by AX and stored in (32-bit) DX:AX
 - There used to be no 32-bit registers
 - If the operand is 32-bit, it is multiplied by EAX and the result is stored in (64-bit) EDX:EAX
 - We don't have 64-bit registers on a 32-bit architecture
- **WARNING:** Multiplication will thus overwrite DX or EDX for 16- or 32-bit values

The imul instruction

- imul, which is used for **signed** numbers, has three formats:

imul src

imul dst, src1

imul dst, src1, src2

- The different combinations are shown in Table 2.2 in the text book
 - Because you are not expected to memorize all options!
- This table uses the typical way in which one specifies operands:
 - reg16: a 16-bit register
 - reg32: a 32-bit register
 - imm8: an 8-bit immediate operand (i.e., a number)
 - mem16: a word of memory
 - etc.
- Let's look at the table

The imul instruction

Will not
overflow
(although the
overflow bit may
be set)

dst	src1	src2	action
	reg/mem8		AX = AL * src1
	reg/mem16		DX:AX = AX * src1
	reg/mem32		EDX:EAX = EAX * src1
reg16	reg/mem16		dst *= src1
reg32	reg/mem32		dst *= src1
reg16	immed8		dst *= immed8
reg32	immed8		dst *= immed8
reg16	immed16		dst *= immed16
reg32	immed32		dst *= immed32
reg16	reg/mem16	immed8	dst = src1*src2
reg32	reg/mem32	immed8	dst = src1*src2
reg16	reg/mem16	immed16	dst = src1*src2
reg32	reg/mem32	immed32	dst = src1*src2

Division

- Two instructions:
 - **div** for unsigned quantities
 - **idiv** for signed quantities
- They perform **integer division**
 - e.g.,: $19 / 4$ produces quotient = 4 remainder = 3
- Only one format for both:
div/idiv src
- If src is an 8-bit quantity:
 - AX is divided by src
 - quotient stored into AL
 - remainder stored into AH
- If src is a 16-bit quantity:
 - DX:AX is divided by src
 - quotient stored into AX
 - remainder stored into DX

Division

- If src is a 32-bit quantity:
 - EDX:EAX is divided by src
 - quotient stored into EAX
 - remainder stored into EDX
- **Warning**: it's very common for programmers to forget initializing DX or EDX before the division
- ANNOYING FEATURE: the argument to div/ idiv must be a register
 - Not a constant
- And yes, this is all cumbersome
- Let's see an example

Division Example

- Say I want to divide 2042 by 13 and I want to have quotient and remainder as 4-byte values
- Here is the code

```
mov      eax, 2042
mov      edx, 0      // important
mov      ecx, 13
idiv     ecx          // must use register
// eax contains 157
// edx contains 1
// since 2042 = 157 * 13 + 1
```

Negation

- There is a convenient instruction to negate an operand: **neg**
- It simply computes the 2's complement of a quantity
- Works on 8-bit, 16-bit, or 32-bit quantities
 - either in registers or in memory
- We'll ignore the content of Section 2.1.5 in the textbook



Example Program in Textbook

- Section 2.1.4 shows a sample program that uses all the arithmetic operations we just saw
- There is nothing particularly difficult about it, especially because overflows are not handled (so the numbers entered had better be “small”)
- Make sure you go through that example and understand how it works
 - You may want to run it as well

High-Level Languages

- Say you have to write a function in C/C++:

```
void f(unsigned int a, unsigned int b) {  
    unsigned int x = a + b;  
    for (unsigned int i=0; i < x; i++) {  
        // do something  
    }  
}
```

- If a user passes numbers whose sum is too big, x will be bogus
- But we cannot check the carry bit
 - We can in assembly, as we'll see
- So we have to check “by hand” :(
- Let's see the code...

High-Level Languages (2)

```
#include <limits.h>

void f(unsigned int a, unsigned int b) {
    if (a > UINT_MAX - b) {
        exit(1); // Overflow
    }

    unsigned int x = a + b;
    for (unsigned int i=0; i < x; i++) {
        // do something
    }
}
```

- You may have had to do this, e.g., when practicing for the coding interview on some sites? (like Leetcode)

High-Level Languages (3)

```
#include <limits.h>

void f(int a, int b) {

    if ((b > 0 ) && (a > INT_MAX - b)) ||
        ((b < 0) && (a < INT_MIN - b)) {
        exit(1);  // Overflow
    }

    int x = a + b;
}
```

- For signed integers, you have to check “both” ends, since you can overflow on either side

High-Level Languages (4)

- As you can see, it's pretty inconvenient, but if you want to write robust code, you have to do it
- What you really want to avoid is a silent overflow
 - Easiest way: always use bigger data types than needed, if possible
- But often, overflow is actually a feature of a program
 - i.e., the program relies on the weird “wrap-around” behavior that happens when you have overflow
- Different languages provide different way of dealing with overflow
- In Java, you can use special “overflow catching” methods of the Math package (e.g., `Math.addExact()`)
- In C/C++ you can give flags to the compiler...

High-Level Languages (4)

- We can ask the C/C++ compiler to add (assembly) code to the check for overflow for all integer operations
 - As we'll learn to do in the next module in assembly
- It's easy for the compiler based on the signed-ness of numbers
 - Insert code to check the carry bit or to check the overflow bit
- If overflow is detected, abort the program
 - But if your program uses overflow as a “feature”, then that will be a problem!
- With gcc: `-ftrapv` will do this for signed overflow
- Alternatively, with gcc: you can do `__builtin_sadd_overflow(a,b,&c)` for an addition that checks overflow and returns true/false
- But that won't work with other compilers....

Do we care?

- Clearly, dealing with overflow is a pain (not in assembly though, as we'll see!)
 - This may be the **one** thing this semester which is better in assembly than in high-level languages
- Let's look at: https://cwe.mitre.org/top25/archive/2020/2020_cwe_top25.html
 - And click at subsequent links
- Many of the examples here says “and then there will be a buffer overflow”
- Stay tuned for a discussion of that vulnerability in the “Subprogram” (post-midterm) module



Conclusion

- One has to be careful when doing arithmetic operations because the processor happily produces results but it's your responsibility to check for overflow/carry bits
- We will have practice quiz on this module **on Thursday**