



Midterm Review

ICS312 Machine-Level and Systems Programming

Henri Casanova (henric@hawaii.edu)

What to Expect?

- **Open note/computer**
- **On Laulima (during lecture period, timed)**
 - A single question with the whole exam as a PDF, a single answer as text
- **Material to review**
 - Practice Quizzes
 - Homework solutions
 - Request solutions via e-mail!
 - Lecture notes
 - Reading assignments in the textbook
 - Especially the examples



Material Covered

- All modules from “Getting Started” to “Bit Operations” (inclusive)

What Questions to Expect

- A few quiz-like questions
 - Study for them by pretending you're teaching the course
 - If you have to go through your notes to answer them all, you'll waste too much time
- Questions like in the homework assignments
 - 2's complement
 - Size modification
 - OF and CF flags after arithmetic operations
 - Memory Layout
 - How does a program modify memory?
 - How does one implement control structures
 - How does one use bitwise operations
- "Write a (small) fragment of assembly code that does"
 - nothing different from what you've done in the homework assignments
- "Here is a program, tell me what it does or fix it"



Numbers

- What's -194 decimal in 2-byte hex?
- What's +36 decimal in 4-byte hex?
- What's F3 in decimal, interpreted as a signed number?
- What's F3 in decimal, interpreted as an unsigned number?

Numbers

- What's -194 in 2-byte hex?
 - $194d = 160d + 34d = 12d * 16d + 2d = 00C2h$
 - flip: FF3D
 - add one: FF3E
- What's +36 in 4-byte hex?
 - $36d = 24h$
- What's F3, interpreted as a signed number?
 - It's negative, so flip: 0C, add one: 0D
 - It's -13d
- What's F3, interpreted as an unsigned number?
 - $15d * 16d + 3d = 243d$



Signed / Unsigned

- The “add” and “sub” operations do the right thing as long as you’re consistent in your interpretations of the operands and of the result
- There is a “mul” instruction for unsigned numbers, and an “imul” instruction for signed numbers
- There is a “div” instruction for unsigned numbers, and an “idiv” instruction for signed numbers

Signed / Unsigned

- The confusing thing:
 - The microprocessor has no notion of whether values are signed or unsigned
 - The programmer has that notion
- If you show me a piece of code written by somebody else, the only way I can tell whether the programmer thinks of numbers as signed or unsigned:
 - Are there any imul? idiv?
 - Does the programmer check to OF flag?
 - Which conditional branch instructions are used?
 - Are there movsx or movzx instructions?

The div instruction

- If src is a 32-bit quantity:
 - EDX:EAX is divided by src
 - quotient stored in EAX
 - remainder stored in EDX
- Don't forget to set EDX to zero!

The imul instruction

Will not
overflow
(although the
overflow bit may
be set)

dst	src1	src2	action
	reg/mem8		AX = AL * src1
	reg/mem16		DX:AX = AX * src1
	reg/mem32		EDX:EAX = EAX * src1
reg16	reg/mem16		dst *= src1
reg32	reg/mem32		dst *= src1
reg16	immed8		dst *= immed8
reg32	immed8		dst *= immed8
reg16	immed16		dst *= immed16
reg32	immed32		dst *= immed32
reg16	reg/mem16	immed8	dst = src1*src2
reg32	reg/mem32	immed8	dst = src1*src2
reg16	reg/mem16	immed16	dst = src1*src2
reg32	reg/mem32	immed32	dst = src1*src2



How to Detect Overflow in Assembly Programs?

UNSIGNED → CARRY FLAG

- jc, jnc

SIGNED → OVERFLOW FLAG

- jo, jno

In both cases we talk of “overflow”!

CF and OF: set or not set?

■ Is CF set?

- Option #1: Do the hex addition and see if you have a left-over carry, in which case $CF=1$
 - That carry would require an additional bit
- Option #2: Reason about the numbers as unsigned and determine whether there will be a carry, in which case $CF=1$, without computing the whole addition

■ Is OF set?

- Think of the two numbers as signed
- If they are of different sign, then $OF=0$ no matter what
- If they are of the same sign, then $OF=1$ only if the sign of the result makes no sense
 - Option #1: Do the hex addition and look at the sign of the result
 - Option #2: Estimate the sign of the result based on magnitudes

A Few Examples

- 1-byte: $AF + 70$
 - Is the Carry Bit set?
 - Is the Overflow Bit set?
- 2-byte: $FF12 + 7FFE$
 - Is the Carry Bit set?
 - Is the Overflow Bit set?
- 1-byte: $AF + 84$
 - Is the Carry Bit set?
 - Is the Overflow Bit set?

A Few Examples

- 1-byte: AF + 70
- Carry flag
 - The “human” method
 - AF is large positive, 70 is large positive, we “overflow”, CF is set
 - The “brute-force” method
 - $AF + 70 = 11F$, CF is set
- Overflow flag
 - AF is negative, 70 is positive, no overflow

A Few Examples

- 2-byte: FF12 + 7FFE
- CF flag
 - The “human” method:
 - We add two huge numbers together: CF is set
 - The “brute-force” method
 - $\text{FF12} + 7\text{FFE} = 17\text{F10}$ and a carry
- OF flag
 - The two numbers are of opposite signs: OF is not set

A Few Examples

- 1-byte: $AF + 84$
- CF Flag
 - The “human” method
 - $A + 8 > F$: carry
 - The “microprocessor” method
 - $AF + 84 = 133$, a carry is generated
- OF flag
 - Both number are negative, so we may have overflow
 - The result is 33, which is positive, so OF is set

Size modification

- When moving a X-byte quantity to a Y-byte quantity, with $X > Y$, you just drop the extra bits on the left
 - May lead to numerically consistent results
 - May lead to numerically inconsistent results
- To increase size one must use
 - Movzx: adds zeros to the left
 - Good to extend the size of unsigned integers
 - Movsx: adds replicas of the sign-bit to the left
 - Good to extend the size of signed integers

Conditional Branches

cmp x, y			
signed		unsigned	
Instruction	branches if	Instruction	branches if
JE	$x = y$	JE	$x = y$
JNE	$x \neq y$	JNE	$x \neq y$
JL, JNGE	$x < y$	JB, JNAE	$x < y$
JLE, JNG	$x \leq y$	JBE, JNA	$x \leq y$
JG, JNLE	$x > y$	JA, JNBE	$x > y$
JGE, JNL	$x \geq y$	JAE, JNB	$x \geq y$

If-then-Else

- The basis of an if-then-else:

cmp XXX

jXX thenblock

; else block

jmp endif

thenblock:

; then block

endif:

Example

- `if ((eax == 0) && (ebx == 1))?`

Example

- if ((eax == 0) && (ebx == 1))?

```
cmp    eax, 0
jnz    elseblock
cmp    ebx, 1
jnz    elseblock
; thenblock
jmp    endif
```

elseblock:

```
; elseblock
```

endif:

Other Example

- if ((eax == 0) || (ebx >= 1))? (signed)

Other Example

- if ((eax == 0) || (ebx >= 1))? (signed)

```
cmp    eax, 0
jz     thenblock
cmp    ebx, 1
jge    thenblock
; elseblock
jmp    endif
```

thenblock:

```
; thenblock
```

endif:

Loops

- Doing: `for (i=3; i<10; i+=2) { body }`

Loops

- Doing: `for (i=3; i<10; i+=2) { body }`

```
                mov ebx, 3
loop1:          ; body
                add ebx, 2
                cmp ebx, 10
                jb  loop1
```

The loop instruction

- Only if
 - You want ecx to be the loop index
 - You want the index to go from some positive value down to zero in increments of 1

```
        mov    ecx, 20
loop1:  body
        loop   loop1
```

The Memory Layout

■ Little Endian

L1	db	"a", "bc"
L2	dd	0AABBCCAAh
L3	times 4	dw -25
L4	db	"d", 0

```
mov    eax, L2
add    eax, 3
mov    word [eax], 23
mov    ebx, L3
mov    ebx, [ebx]
movsx  eax, bh
mov    [L3], eax
```

The Memory Layout

■ Little Endian

L1	db	"a","bc"
L2	dd	0AABBCCAAh
L3	times 4	dw -25
L4	db	"d", 0

```
mov     eax, L2
add     eax, 3
mov     word [eax], 23
mov     ebx, L3
mov     ebx, [ebx]
movsx   eax, bh
mov     [L3], eax
```

61 62 63 AA CC BB AA E7 FF E7 FF E7 FF E7 FF 64 00

The Memory Layout

■ Little Endian

L1	db	"a","bc"
L2	dd	0AABBCCAAh
L3	times 4	dw -25
L4	db	"d", 0

```
mov     eax, L2
add     eax, 3
mov     word [eax], 23
mov     ebx, L3
mov     ebx, [ebx]
movsx   eax, bh
mov     [L3], eax
```

61 62 63 AA CC BB AA E7 FF E7 FF E7 FF E7 FF 64 00



The Memory Layout

■ Little Endian

L1	db	"a","bc"
L2	dd	0AABBCCAAh
L3	times 4	dw -25
L4	db	"d", 0

```
mov    eax, L2
add    eax, 3
mov    word [eax], 23
mov    ebx, L3
mov    ebx, [ebx]
movsx  eax, bh
mov    [L3], eax
```

61 62 63 AA CC BB AA E7 FF E7 FF E7 FF E7 FF 64 00



The Memory Layout

■ Little Endian

L1	db	"a","bc"
L2	dd	0AABBCCAAh
L3	times 4	dw -25
L4	db	"d", 0

```
mov     eax, L2
add     eax, 3
mov     word [eax], 23
mov     ebx, L3
mov     ebx, [ebx]
movsx   eax, bh
mov     [L3], eax
```

61 62 63 AA CC BB 17 00 FF E7 FF E7 FF E7 FF 64 00



The Memory Layout

■ Little Endian

L1	db	"a", "bc"
L2	dd	0AABBCCAAh
L3	times 4	dw -25
L4	db	"d", 0

```
mov     eax, L2
add     eax, 3
mov     word [eax], 23
mov    ebx, L3
mov     ebx, [ebx]
movsx   eax, bh
mov     [L3], eax
```

61 62 63 AA CC BB 17 00 FF E7 FF E7 FF E7 FF 64 00



The Memory Layout

■ Little Endian

L1	db	"a","bc"
L2	dd	0AABBCCAAh
L3	times 4	dw -25
L4	db	"d", 0

mov	eax, L2
add	eax, 3
mov	word [eax], 23
mov	ebx, L3
mov	ebx, [ebx]
movsx	eax, bh
mov	[L3], eax

61 62 63 AA CC BB 17 00 FF E7 FF E7 FF E7 FF 64 00

ebx = FF E7 FF 00

The Memory Layout

■ Little Endian

L1	db	"a", "bc"
L2	dd	0AABBCCAAh
L3	times 4	dw -25
L4	db	"d", 0

```
mov     eax, L2
add     eax, 3
mov     word [eax], 23
mov     ebx, L3
mov     ebx, [ebx]
movsx   eax, bh
mov     [L3], eax
```

61 62 63 AA CC BB 17 00 FF E7 FF E7 FF E7 FF 64 00

ebx = FF E7 FF 00

eax = FF FF FF FF

The Memory Layout

■ Little Endian

L1	db	"a", "bc"
L2	dd	0AABBCCAAh
L3	times 4	dw -25
L4	db	"d", 0

```
mov     eax, L2
add     eax, 3
mov     word [eax], 23
mov     ebx, L3
mov     ebx, [ebx]
movsx   eax, bh
mov     [L3], eax
```

61 62 63 AA CC BB 17 FF FF FF FF E7 FF E7 FF 64 00

ebx = FF E7 FF 00
eax = FF FF FF FF

Bit Operations

- What does this code print out?

```
mov     ax, 0035Dh
sal     ah, 6
sar     ax, 2
shr     ah, 2
xor     ah, al
not     ah
sar     ah, 4
movsx   eax, ah
call    print_int
```

Bit Operations

EAX

```
mov    ax, 0035Dh
sal     ah, 6
sar     ax, 2
shr     ah, 2
xor     ah, al
not     ah
sar     ah, 4
movsx   eax, ah
call    print_int
```

????????	????????	00000011	01011101
????????	????????	11000000	01011101
????????	????????	11110000	00010111
????????	????????	00111100	00010111
????????	????????	00101011	00010111
????????	????????	11010100	00010111
????????	????????	11111101	00010111
11111111	11111111	11111111	11111101

FFFFFFFFD

flip: 00000002

+1: 00000003

the code prints out: -3



Counting bits

- The only real example we've seen is how to count bits (set to 1) in some register
- Should we look at this again?

The Carry bits and Shifts

- Remember that when you do a shift, the last bit shifted out ends up in the carry bit
- Remember that the **adc** instruction is very useful as it adds the carry bit to a register:
 `adc eax, 12 ; eax += 12 + carry`
 `adc al, 0 ; al += 0 + carry`
- Let's review a few simple things with bitmasks...

Example #1

- Code to flip the n^{th} bit of EAX, counting from right to left from the rightmost bit, where n is stored in cl

Example #1

- Code to flip the n^{th} bit of EAX, counting from right to left from the rightmost bit, where n is stored in cl

```
mov    ebx, 1
```

```
shl    ebx, cl    ; only cl works here
```

```
xor    eax, ebx
```

Example #2

- Create in `eax` a 32-bit bitmask that looks like 0's followed by `n` 1's, where `n` is stored in `cl`

Example #2

- Create in `eax` a 32-bit bitmask that looks like 0's followed `n` 1's, where `n` is stored in `cl`

```
mov     eax, 0FFFFFFFFh
shl     eax, cl
not     eax
```

Example #3

- Create in `eax` a 32-bit bitmask that contains n 1's, followed by $32-2n$ 0's, followed by n 1's, where n is stored in `cl`

Example #3

- Create in `eax` a 32-bit bitmask that contains n 1's, followed by $32-2n$ 0's, followed by n 1's, where n is stored in `cl`

```
mov     eax, 0FFFFFFFFh
shl     cl, 1 ; multiply cl by 2
shl     eax, cl
shr     cl, 1 ; divide cl by 2
ror     eax, cl
not     eax
```



Mystery Program

- In close-notes exams, I often have “mystery program” questions
- I don’t do this for open-notes/open-computer exams
- But it may still good practice for you to see if you can do the following mystery program problems...

Mystery Program Example

L resw 10

. . .

mov ebx, L

add ebx, 18

mov ecx, 0

loop1: movsx word eax, [ebx]

add ecx, eax

sub ebx, 2

cmp ebx, L

jnz loop1

Mystery Program Example

L resw 10

...

mov ebx, L

add ebx, 18

mov ecx, 0

loop1: movsx word eax, [ebx]

add ecx, eax

sub ebx, 2

cmp ebx, L

jnz loop1

computes the sum of the
elements in an array of
10 2-byte values, which
starts at address L

Yet another mystery program..

```
segment .data
    msg      db "Enter an integer: ",
0
    success db "Success", 0
```

```
segment .text
```

```
again:  . . .
        mov     eax, msg
        call    print_string
        call    read_int
        and     al, 0
        cmp     eax, 0
        jnz     again

        mov     eax, success
        call    print_string
        call    print_nl
        . . .
```

Yet another mystery program..

```
segment .data
    msg      db "Enter an integer: ",
0
    success db "Success", 0
```

```
segment .text
```

```
again: . . .
    mov     eax, msg
    call    print_string
    call    read_int
    and     al, 0
    cmp     eax, 0
    jnz     again

    mov     eax, success
    call    print_string
    call    print_nl
    . . .
```

repeatedly asks the user for an integer until that integer is between 0 and 255, then prints "Success".

Other Mystery Program

```
segment .bss
    L      resb    10
```

```
segment .text
```

```
    . . .
    mov     ebx, L
    add     ebx, 9
```

```
    mov     ecx, 0
```

```
for:
```

```
    mov     dl, [ebx]
    shr     dl, 1
    jc      nope
    inc     ecx
```

```
nope:
```

```
    dec     ebx
    cmp     ebx, L
    jnz     for
```

```
    mov     eax, ecx
    call    print_int
    call    print_nl
```

```
    . . .
```

Other Mystery Program

```
segment .bss
    L      resb    10
```

```
segment .text
```

```
    . . .
```

```
    mov     ebx, L
```

```
    add     ebx, 9
```

```
    mov     ecx, 0
```

```
for:
```

```
    mov     dl, [ebx]
```

```
    shr     dl, 1
```

```
    jc      nope
```

```
    inc     ecx
```

```
nope:
```

```
    dec     ebx
```

```
    cmp     ebx, L
```

```
    jnz     for
```

```
    mov     eax, ecx
```

```
    call    print_int
```

```
    call    print_nl
```

```
    . . .
```

prints the number of
even values among the
first 10 1-byte values at
address L

Yet another mystery program..

```
segment .bss
    L        resd 10

segment .text
    . . .
    mov      ecx, 0
    mov      edx, 10
    mov      ebx, L
for:
    mov      eax, [ebx]
    shr      eax, 1
    jc       no
    shl      eax, 1
    cmp      eax, ecx
    jle      no
    mov      ecx, eax
no:
    add      ebx, 4
    dec      edx
    jnz      for

    mov      eax, ecx
    call     print_int
    call     print_nl
    . . .
```

Yet another mystery program..

```
segment .bss
    L      resd 10

segment .text
    . . .
    mov     ecx, 0
    mov     edx, 10
    mov     ebx, L
for:
    mov     eax, [ebx]
    shr     eax, 1
    jc      no
    shl     eax, 1
    cmp     eax, ecx
    jle     no
    mov     ecx, eax
no:
    add     ebx, 4
    dec     edx
    jnz     for

    mov     eax, ecx
    call    print_int
    call    print_nl
    . . .
```

prints the largest **even**
values among the first 10
4-byte values at address
L



Questions....