



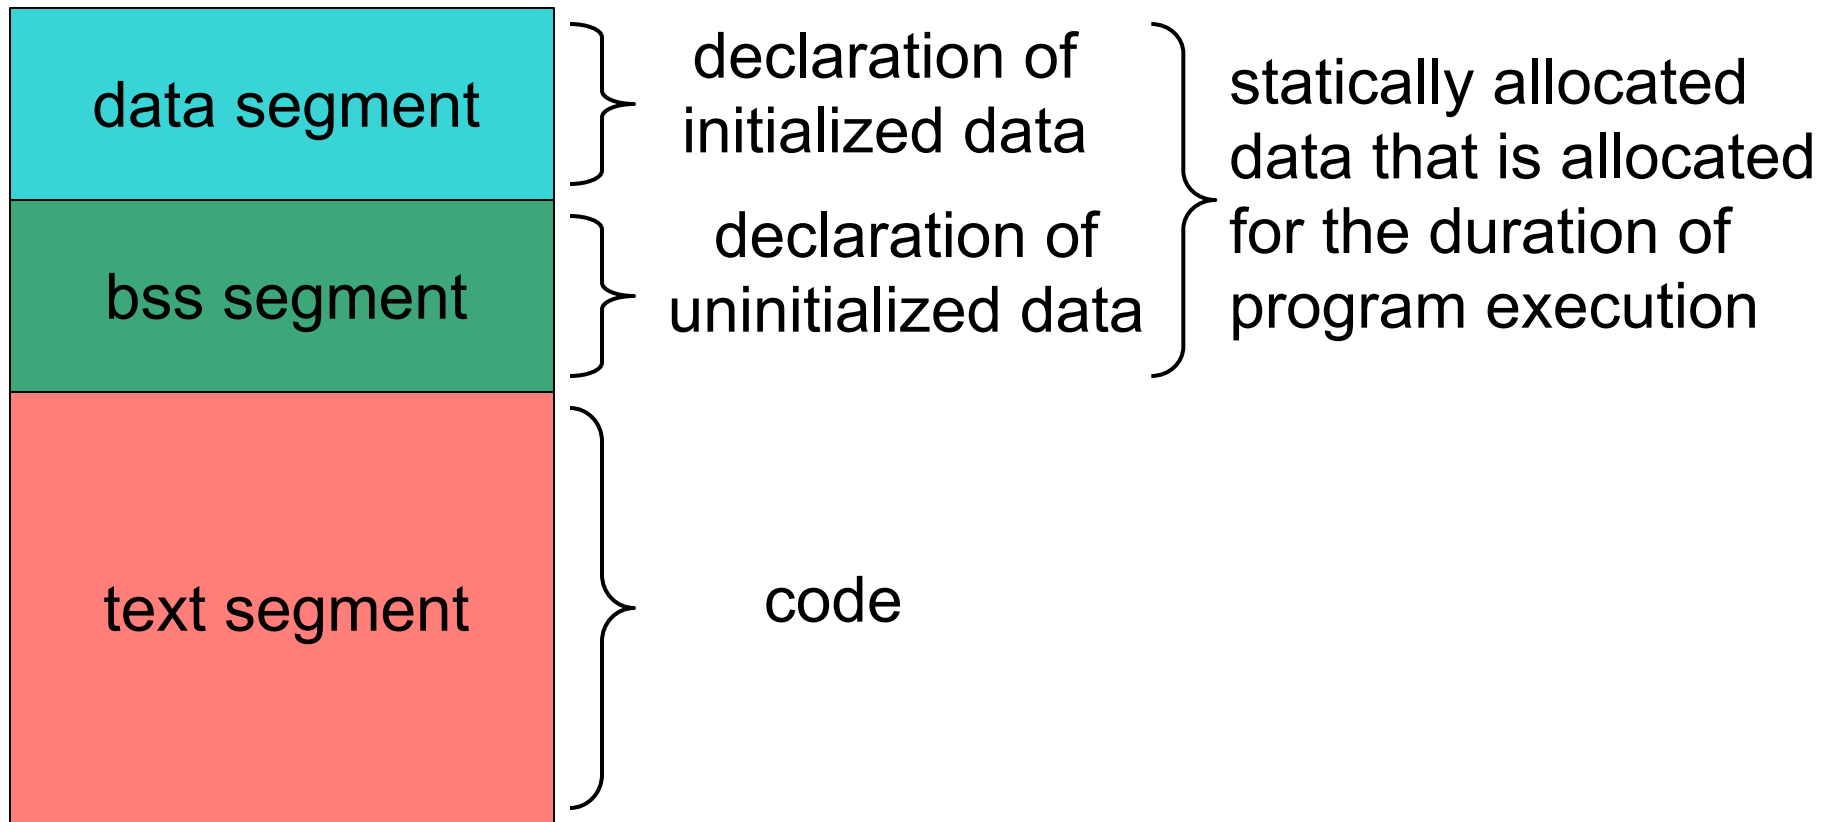
NASM: data and bss

Introduction

ICS312
Machine-Level and
Systems Programming

Henri Casanova (henric@hawaii.edu)

NASM Program Structure



The data and bss segments

- Both segments contains **data directives** that **declare** pre-allocated zones of memory
- There are two kinds of data directives
 - **DX directives:** **initialized** data (D = “defined”)
 - **RESX directives:** **uninitialized** data (RES = “reserved”)
- The “X” above refers to the data size:

Unit	Letter(X)	Size in bytes
byte	B	1
word	W	2
double word	D	4
quad word	Q	8
ten bytes	T	10

The DX data directives

- One declares a zone of initialized memory using three elements:
 - **Label**: the name used in the program to refer to that zone of memory
 - A pointer to the zone of memory, i.e., an address
 - **DX**, where X is the appropriate letter for the size of the data being declared
 - **Initial value**, with encoding information
 - default: decimal
 - b: binary
 - h: hexadecimal
 - o: octal
 - quoted: ASCII

DX Examples

- L1 db 0
 - 1 byte, **whose address is named L1**, initialized to 0
- Henri dw 1000
 - 2-byte word, named Henri, initialized to 1000
- L3 db 110101b
 - 1 byte, named L3, initialized to 110101 in binary
- what db 0A2h
 - 1 byte, named what, initialized to A2 in hex (**note the '0'**)
- L5 db 17o
 - 1 byte, named L5, initialized to 17 in octal ($1*8+7=15$ in decimal)
- L6 dd 0FFFF1A92h (**note the '0'**)
 - 4-byte double word, named L6, initialized to FFFF1A92 in hex
- L7 db "A"
 - 1 byte, named L7, initialized to the ASCII code for "A" (65d)

ASCII Code

- Associates 1-byte numerical codes to characters
 - Unicode, proposed much later, uses 2 bytes and thus can encode 2^8 times more characters (room for all languages, Chinese, Japanese, accents, etc.)
- A few values to know:
 - 'A' is 65d / 41h
 - 'B' is 66d / 42h, etc...
 - 'a' is 97d / 61h
 - 'b' is 98d / 62h, etc...

DX for multiple elements

- L8 db 0, 1, 2, 3
 - Defines 4 1-byte values, initialized to 0, 1, 2 and 3
 - L8 is a pointer to (i.e., the address of) the first byte
- The above is equivalent (in terms of memory content) to:
 - L8 db 0
 - L9 db 1
 - L10 db 2
 - L11 db 3
- The only difference is that in the second version we have a name (label) for the address of each of the four bytes

Strings as sequences of chars

- L9 db “w”, “o”, ‘r’, ‘d’, 0
 - Defines 5 1-byte values, the first 4 being initialized by an ASCII code (i.e., a character)
 - Defines a **null-terminated** string, initialized to “word\0”
 - L9 is a pointer to the beginning of the string (i.e., the address of the first character of the string)

- L10 db “word”, 0
 - Equivalent to the above, more convenient to write

DX with the times qualifier


- Say you want to declare 100 bytes all initialized to 0
- NASM provides a nice shortcut to do this, the “times” qualifier
- `L11 times 100 db 0`
 - Equivalent to `L11 db 0,0,0,....,0` (100 times)

Uninitialized Data

- The RESX directive is very similar to the DX directive, but **always specifies the number of memory elements**
- L20 resw 100
 - 100 uninitialized 2-byte values
 - L20 is a pointer to the first 2-byte value
- stuff resb 1
 - 1 uninitialized byte named stuff

Data segment example

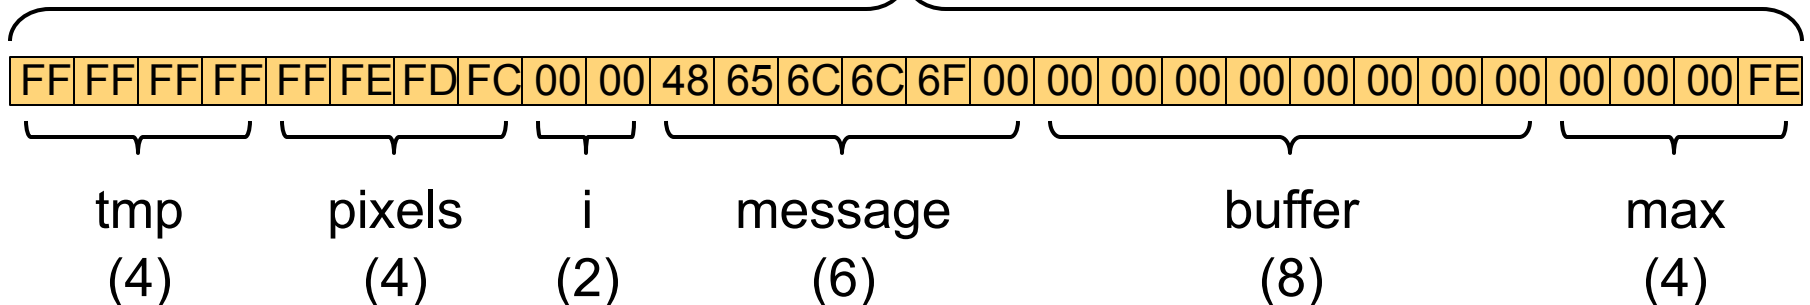
tmp	dd	-1
pixels	db	0FFh, 0FEh, 0FDh, 0FCh
i	dw	0
message	db	"H", "e", "llo", 0
buffer	times 8	db 0
max	dd	254



Data segment example

tmp	dd	-1
pixels	db	0FFh, 0FEh, 0FDh, 0FCh
i	dw	0
message	db	"H", "e", "llo", 0
buffer	times 8	db 0
max	dd	254

28 bytes



Endianness?

max	dd	254
-----	----	-----

00	00	00	FE
----	----	----	----

max

- In the previous slide I showed the above 4-byte **memory** content for a double-word that contains $254 = 000000FEh$
- While this seems to make sense, it turns out that Intel processors do not do this!
 - Yes, the last 4 bytes shown in the previous slide are wrong
- The scheme shown above (i.e., bytes in memory follow the “natural” order): **Big Endian**
- Instead, Intel processors use **Little Endian**:

FE	00	00	00
----	----	----	----


max

Little Endian


```
mov eax, 0AFBBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```

Registers

eax




ebx



Memory

[M1]



Little Endian

```
mov eax, 0AFBBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```

Registers

eax

AF	BB	CC	DD
----	----	----	----

ebx

--	--	--	--

Memory

[M1]

--	--	--	--

Little Endian

```
mov eax, 0AFBBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```

Registers

eax

AF	BB	CC	DD
----	----	----	----

ebx

--	--	--	--

Memory

[M1]

DD	CC	BB	AF
----	----	----	----

Little Endian

```
mov eax, 0AFBBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```

Registers

eax

AF	BB	CC	DD
----	----	----	----

ebx

AF	BB	CC	DD
----	----	----	----

Memory

[M1]

DD	CC	BB	AF
----	----	----	----

Little Endian

```
mov eax, 0AFBBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```



In-register byte order and in-memory byte order, within a single multi-byte value, are reversed!

Little/Big Endian

- Motorola and IBM processors use(d) Big Endian
- Intel/AMD uses Little Endian (used in this class)
- When writing code in a high-level language one rarely cares
 - Although in C one can definitely expose the Endianness of the computer
 - And thus one can write C code that's not portable between an IBM and an Intel!!!
- This only matters when writing **multi-byte** quantities to memory and reading them differently (e.g., byte per byte)
- When writing assembly code one often does not care, but we'll see several examples when it matters, so it's important to know this *inside out*
- Some processors are configurable (either in hardware or in software) to use either type of endianness (e.g., MIPS processor)

Example

pixels	times 4	db	0FDh
x	dd	00010111001101100001010111010011b	
blurb	db	"ad", "b", "h", 0	
buffer	times 10	db	14o
min	dw	-19	

- What is the layout and the content of the data memory segment on a Little Endian machine?
 - Byte per byte, in hex

Example

pixels	times 4	db	0FDh
x	dd	00010111001101100001010111010011b	
blurb	db	"ad", "b", "h", 0	
buffer	times 10	db	14o
min	dw	-19	

- First thing to do: identify the multi-byte quantities

Example


pixels	times 4	db	0FDh
x	dd	00010111001101100001010111010011b	
blurb	db	"ad", "b", "h", 0	
buffer	times 10	db	14o
min	dw	-19	

- First thing to do: identify the multi-byte quantities
 - EVERYTHING THAT'S NOT DECLARED AS "db" IS MULTI-BYTE
 - (L db "stuff" is NOT MULTI-BYTE)
- In the above: **x** and **min** above are multi-byte values

Example


pixels	times	4	db	0FDh
x	dd	00010111001101100001010111010011b		
blurb	db	"ad", "b", "h", 0		
buffer	times	10	db	14o
min	dw	-19		

25 bytes



Example

pixels	times	4	db	0FDh
x	dd	00010111001101100001010111010011b		
blurb	db	"ad", "b", "h", 0		
buffer	times	10	db	14o
min	dw	-19		



Note that bits with each byte are NOT reversed



What we'll do in Class

- We'll do some “what are the bytes in RAM given the data segment specification?” practice
- Then I'll keep lecturing about this, which will lead to more in-class practices
- Make sure you really have your terminology down pat before you come to class
 - The ‘b’, ‘w’, ‘d’ things
 - Syntax of the data segment declarations, including the “times” qualifier
 - Little-Endian weirdness